


Arithmetic Logic Unit

Part 5

1



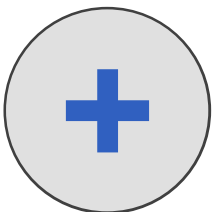
Adding Binary Integers

$1 + 1 = 10$

2

Adding Binary Integers

- Computer's add binary numbers the same way that we do with decimal
- Columns are aligned, added, and "1's" are carried to the next column
- In computer processors, this component is called an *adder*



Fall 2024 Semester 2Bn - Cnbn - CSnB 2B 3

3

Adding Base 10 Numbers

	1		1				
	2	7	8	1			
+	3	7	2	1			
	<hr/>						
	6	5	0	2			

Fall 2024 Semester 2Bn - Cnbn - CSnB 2B 4

4

Adding Binary Example

118

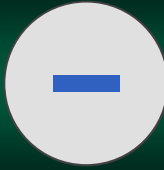
51

169

	1		1		1		1		1			
	0	1	1	1	0	1	1	0				
+	0	0	1	1	0	0	1	1				
	<hr/>											
	1	0	1	0	1	0	0	1				

Fall 2024 Semester 2Bn - Cnbn - CSnB 2B 5

5



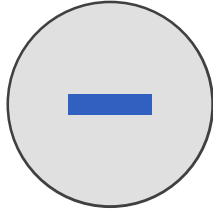
Negative Binary Integers

Have a positive attitude about negatives

6

Negative Binary Numbers

- When we write a negative number, we generally use a "-" as a prefix character
- However, binary numbers can only store ones and zeros



Fall 2024

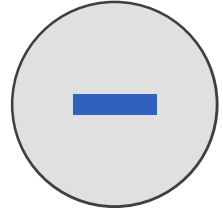
Secrets: State - Cook - CSU 35

7

7

Negative Binary Numbers

- So, how we store a negative a number?
- When a number can represent both positive and negative numbers, it is called a *signed integer*
- Otherwise, it is *unsigned*



Fall 2024

Secrets: State - Cook - CSU 35

8

8

Signed Magnitude

- One approach is to use the most significant bit (msb) to represent the negative sign
- If positive, this bit will be a zero
- If negative, this bit will be a 1
- This gives a byte a range of -127 to 127 rather than 0 to 255

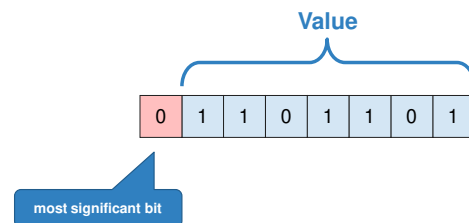
Fall 2024

Secrets: State - Cook - CSU 35

9

9

Signed Magnitude



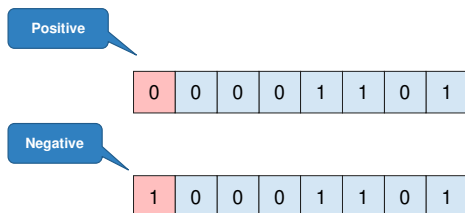
Fall 2024

Secrets: State - Cook - CSU 35

10

10

Signed Magnitude: 13 and -13



Fall 2024

Secrets: State - Cook - CSU 35

11

11

Signed Magnitude Drawback #1

- When two numbers are added, the system needs to check and sign bits and act accordingly
- For example:
 - if both numbers are positive, add values
 - if one is negative subtract it from the other
 - etc...
- There are also rules for subtracting

Fall 2024

Secrets: State - Cook - CSU 35

12

12

Signed Magnitude Drawback #2

- Also, signed magnitude also can store a positive and negative version of zero
- Yes, there are two zeroes!
- Imagine having to write Java code like...

```
if (x == +0 || x == -0)
```

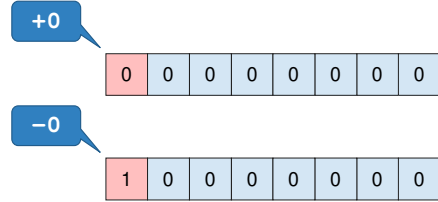
Fall 2024

Segment 8: Bits - CS50.10

13

13

Oh noes! Two zeros?



Fall 2024

Segment 8: Bits - CS50.10

14

14

1's Complement

- Rather than use a sign bit, the value can be made negative by *inverting* each bit
 - each 1 becomes a 0
 - each 0 becomes a 1
- Result is a "complement" of the original
- This is logically the same as subtracting the number from 0

Fall 2024

Segment 8: Bits - CS50.10

15

15

Advantages / Disadvantages

- Advantages over signed magnitude
 - very simple rules for adding/subtracting
 - numbers are simply added:
5 - 3 is the same as 5 + -3
- Disadvantages
 - positive and negative zeros still exist
 - so, it's not a perfect solution

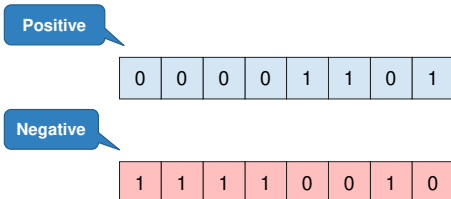
Fall 2024

Segment 8: Bits - CS50.10

16

16

1's Complement: 13 and -13



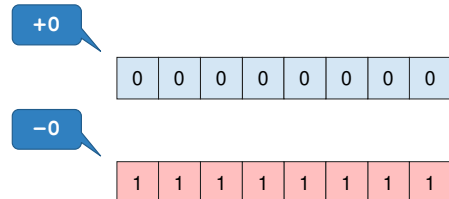
Fall 2024

Segment 8: Bits - CS50.10

17

17

1's Complement Has Two Zeros



Fall 2024

Segment 8: Bits - CS50.10

18

18

2's Complement

- Practically all computers use *2's Complement*
- Similar to 1's complement, but after the number is inverted, 1 is added to the result
- Logically the same as:
 - subtracting the number from 2^n
 - where n is the total number of bits in the integer



Fall 2024

Secureworks, State - Cook - CSIS 35

19

19

2's Complement Advantages

- Since negatives are subtracted from 2^n
 - they can simply be added
 - the extra carry 1 (if it exists) is discarded
 - this simplifies the hardware considerably since the processor only has to add
- The +1 for negative numbers...
 - makes it so there is only one zero
 - values range from -128 to 127

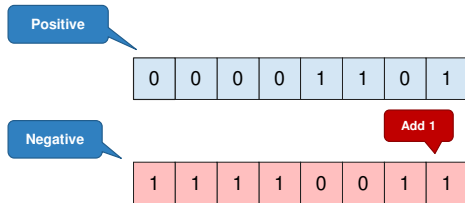
Fall 2024

Secureworks, State - Cook - CSIS 35

20

20

2's Complement: 13 and -13



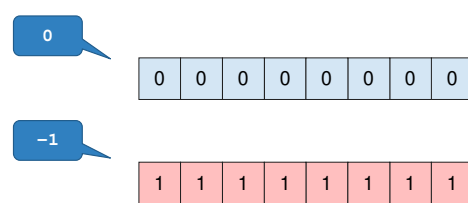
Fall 2024

Secureworks, State - Cook - CSIS 35

21

21

Just One Zero!



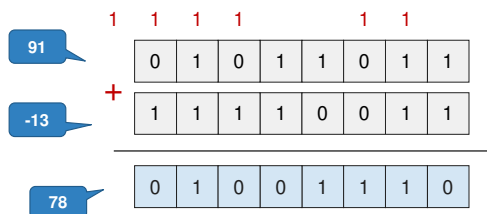
Fall 2024

Secureworks, State - Cook - CSIS 35

22

22

Adding 2's Complement



Fall 2024

Secureworks, State - Cook - CSIS 35

23

23

Unsigned or Signed?

- In reality, processors don't know (*or care*) if a number is unsigned or signed
- The hardware works the same either way
- It's your responsibility to keep track if it's signed/unsigned



Fall 2024

Secureworks, State - Cook - CSIS 35

24

24

It's Your Responsibility




- In many cases, you must use the correct instruction - based on what *you* are treating as a constant or variable
- In great programming, never comes great responsibility



Fall 2024 Semester 2B - CS4 - CSU 25 25

25




Relative Addressing

Jumping of the instruction pointer

26

Relative Addressing

- In *relative addressing*, a value is added to a instruction pointer (e.g. program counter)
- This allows access a fixed number of bytes *up or down* from the instruction pointer



Fall 2024 Semester 2B - CS4 - CSU 25 27

27

Relative Addressing

- Often used in conditional jump statements
 - jumps are often short – not a large number of instructions
 - so, the instruction only stores the value to add to the program counter
 - practically all processors use this approach
- Also used to access local data – load/store

Fall 2024 Semester 2B - CS4 - CSU 25 28


28

Relative Addressing Advantages

- The instruction can just store the *difference* (in bytes) from the current instruction address
- It takes less storage than a full 64-bit address
- It also allows a program to be stored anywhere in memory – *and it will still work!*

Fall 2024 Semester 2B - CS4 - CSU 25 29

29



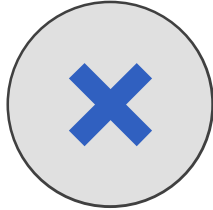
Multiplying Binary Numbers

$11 \times 11 = 1001$

30

Multiplying Binary Numbers

- Many processors today provide complex mathematical instructions
- However, the processor only needs to know how to add
- Historically, multiplication was performed with successive additions



Fall 2024

Secrets of the Machine - CS50.02

31

31

Multiplying Scenario

- Let's say we have two variables: **A** and **B**
- Both contain integers that we need to multiply
- Our processor can only add (and subtract using 2's complement)
- How do we multiply the values?

Fall 2024

Secrets of the Machine - CS50.02

32

32

Multiplying: The Bad Way



- One way of multiplying the values is to create a For Loop using one of the variables – **A** or **B**
- Then, inside the loop, continuously add the other variable to a running total

Fall 2024

Secrets of the Machine - CS50.02

33

33

Multiplying: The Bad Way

```
total = 0;
for (i = 0; i < A; i++)
{
    total += B;
}
```

Fall 2024

Secrets of the Machine - CS50.02

34

34

Multiplying: The Bad Way

- If **A** or **B** is large, then it could take a long time
- This is incredibly inefficient
- Also, given that **A** and **B** could contain drastically different values – the number of iterations would vary
- Required time is not constant



Fall 2024

Secrets of the Machine - CS50.02

35

35

Multiplying: The Best Way



- Computers can multiply by using long multiplication – *just like you do*
- Number of additions is fixed to 8, 16, 32, 64 depending on the size of the integer
- The following example multiplies 2 unsigned 4-bit numbers

Fall 2024

Secrets of the Machine - CS50.02

36

36

Unsigned Integer: 13 × 10

1101

×

1010

0000

+

Fall 2024

Segment 8: Bits - CS50.06

37

37

Unsigned Integer: 13 × 10

1101

×

1010

0000

1101

+

Fall 2024

Segment 8: Bits - CS50.06

38

38

Unsigned Integer: 13 × 10

1101

×

1010

0000

1101

0000

+

Fall 2024

Segment 8: Bits - CS50.06

39

39

Unsigned Integer: 13 × 10

1101

×

1010

0000

1101

0000

1101

+

Fall 2024

Segment 8: Bits - CS50.06

40

40

Unsigned Integer: 13 × 10

1101

×

1010

0000

1101

0000

1101

10000010

+

130

Fall 2024

Segment 8: Bits - CS50.06

41

41

Multiplication Doubles the Bit-Count

When two numbers are multiplied, the product will have twice the number of digits

Examples:

8-bit × 8-bit → 16-bit

16-bit × 16-bit → 32-bit

32-bit × 32-bit → 64-bit

64-bit × 64-bit → 128-bit

Fall 2024

Segment 8: Bits - CS50.06

42

42

7

Multiplication Doubles the Bit-Count

- So, how do we store the result?
- It is often too large to fit into any single existing register
- Processors can...
 - fit the result in the original bit-size *(and raise an overflow if it does not fit)*
 - ...or store the new double-sized number

Fall 2024

Sec 10.10: x86-64 - CS: 10

43

43



x86 Mathematics

Complex Math is Complex

44

Add & Subtract

- The Add and Subtract instructions take two operands and store the result in the first operand
- This is the same as the `+=` and `-=` operators used in Visual Basic .NET, C, C++, Java, etc...



Fall 2024

Sec 10.10: x86-64 - CS: 10

45

45

Addition

ADD *target*, *value*

Immediate, Register, Memory

Register, Memory

Fall 2024

Sec 10.10: x86-64 - CS: 10

46

46

Subtraction

SUB *target*, *value*

Immediate, Register, Memory

Register, Memory

Fall 2024

Sec 10.10: x86-64 - CS: 10

47

47

Negate (2's complement)

NEG *register*

Fall 2024

Sec 10.10: x86-64 - CS: 10

48

48

Example: Simple Add

```
MOV rax, 17
ADD rax, 2
```

Move value into RAX

RAX += 2

49



x86 Multiplication

Complex Math is Complex

50

Multiplication & Division

- The x86 treats multiplication quite differently than add/subtract
- Why? Intel was designed as a business processor and high-precision math is paramount



51

Multiplication Review

- Remember: when two n bit numbers are multiplied, result will be $2n$ bits
- So...
 - two 8-bit numbers \rightarrow 16-bit
 - two 16-bit numbers \rightarrow 32-bit
 - two 32-bit numbers \rightarrow 64-bit
 - two 64-bit numbers \rightarrow 128-bit



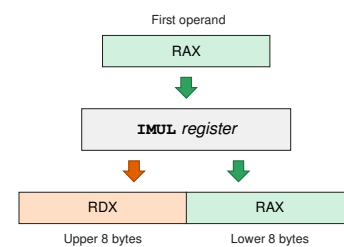
52

Multiplication on the x86

- Intel stores the product into two registers
 - RAX will contain the lower 8 bytes
 - RDX will contain the upper 8 bytes
- This maintains the high-precision result
- Instruction inputs are strange
 - first operand is must be stored in RAX
 - second operand must be a register or memory

53

x86 Multiplication



54

Multiply - Signed

IMUL *operand*

Register or Memory only

55

Multiply - Unsigned

MUL *operand*

Register or Memory only

56

Signed Multiply: 1846 by 42

```
MOV    rax, 1846    #First operand
MOV    rbx, 42      #Need register for MUL
IMUL    rbx          #RAX gets low 8 bytes
                        #RDX gets high 8 bytes
```

57

Multiplication Tips

- Even though you are just using RAX as input, both RAX and RDX will change
- Be aware that you might lose important data, and backup to memory if needed



58

Additional x86 Multiply Instructions

- Over time, designers requested a low-precision version of multiplication
- Intel added "short" IMUL instructions that store into a single register
- Please Note: these do not exist for MUL



59

IMUL (few more combos)

IMUL *target, value*

Immediate, Register, Memory

Register

60

Signed Multiply: 1846 by 42

```
MOV    rax, 1846
IMUL   rax, 42
```

This works, but could cause an overflow

61



Extending Byte Size

Converting from 8-bit to 16-bit and more

62

Extending Unsigned Integers

- Often in programs, data needs to be moved to an integer with a larger number of bits
- For example, an 8-bit number is moved to a 16-bit representation



63

Extending Unsigned Integers

- For unsigned numbers is fairly easy – just add zeros to the left of the number
- This, naturally, is how our number system works anyway: $456 = 000456$



64

Unsigned 13 Extended

0 0 0 0 1 1 0 1



0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1

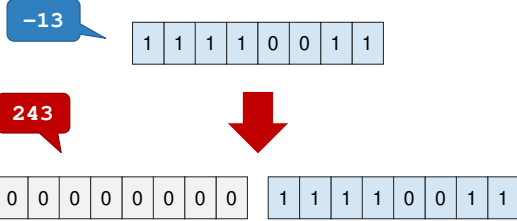
65

Extending Signed Integers

- When the data is stored in a signed integer, the conversion is a little more complex
- Simply adding zeroes to the left, will *convert a negative value to a positive one*
- Each type of signed representation has its own set of rules

66

2's Complement Incorrectly Done



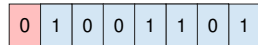
67

Sign Magnitude Extension

- In signed magnitude, the most-significant bit (msb) stores the negative sign
- The new sign-bit needs to have this value
- Rules:
 - copy the old sign-bit to the new sign-bit
 - fill in the rest of the new bits with zeroes – *including the old sign bit*

68

Sign Magnitude Extended: +77



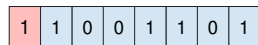
69

Sign Magnitude Extended: +77



70

Sign Magnitude Extended: -77



71

Sign Magnitude Extended: -77



72

1's & 2's Complement Extension

- 1's and 2's Complement is very simple to convert to a larger representation
- Remember that we inverted the bits and added 1 to get a negative value
- Rule: copy the old most-significant bit to all the new bits



Fall 2024

Segment 8: Bits - C&C - CSU 35

73

73

1's & 2's Complement Extended: +77

0	1	0	0	1	1	0	1
---	---	---	---	---	---	---	---

Fall 2024

Segment 8: Bits - C&C - CSU 35

74

74

1's & 2's Complement Extended: +77

0	0	0	0	0	0	0	0
0	1	0	0	1	1	0	1

Fall 2024

Segment 8: Bits - C&C - CSU 35

75

75

1's & 2's Complement Extended: -77

1	0	1	1	0	0	1	1
---	---	---	---	---	---	---	---

Fall 2024

Segment 8: Bits - C&C - CSU 35

76

76

1's & 2's Complement Extended: -77

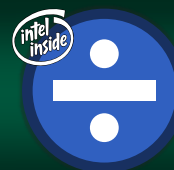
1	1	1	1	1	1	1	1
1	0	1	1	0	0	1	1

Fall 2024

Segment 8: Bits - C&C - CSU 35

77

77



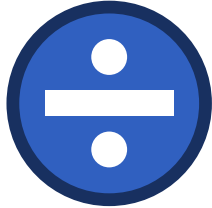
x86 Division

Complex Math is Complex

78

Division on the x86

- Division on the x86 is very interesting
- Since multiplication stores into two registers, divide uses these as the numerator
 - RAX** contains the lower 8 bytes
 - RDX** contains the upper 8 bytes



Fall 2024

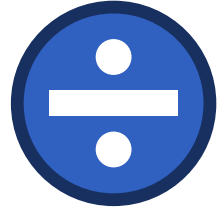
Secretworks, Siles - Cook - CSU 35

79

79

Division on the x86

- These two registers are also used for the result
 - RAX** will contain the quotient (the whole number)
 - RDX** will contain the remainder



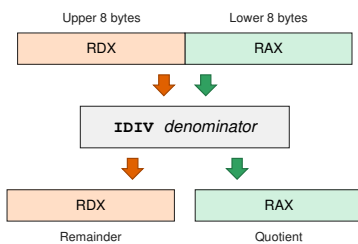
Fall 2024

Secretworks, Siles - Cook - CSU 35

80

80

x86 Division



Fall 2024

Secretworks, Siles - Cook - CSU 35

81

81

Divide - Signed

IDIV *denominator*

Register or Memory only

Fall 2024

Secretworks, Siles - Cook - CSU 35

82

82

Divide - Unsigned

DIV *denominator*

Register or Memory only

Fall 2024

Secretworks, Siles - Cook - CSU 35

83

83

Dividing Rules

- The numerator **must** be expanded to the destination size (twice the original)
- Why? Multiplication doubles the number of digits; division does the opposite
- This must be done **before** the division - otherwise the result will be incorrect

Fall 2024

Secretworks, Siles - Cook - CSU 35

84

84

On the Intel...

- You **must** setup RDX **before** you divide
- For unsigned: store 0 into it
- For signed-division:
 - RAX needs must be *sign-extended* into RDX
 - there are special instructions



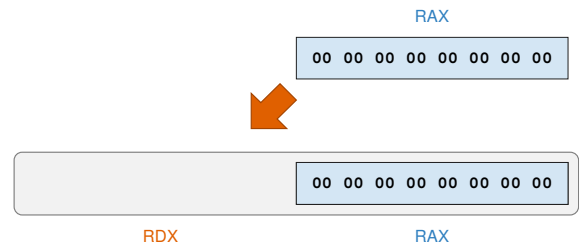
Fall 2024

Secureware State - Cook - CSU 35

85

85

Sign Extend Example



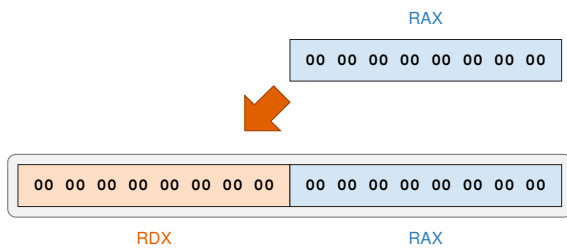
Fall 2024

Secureware State - Cook - CSU 35

86

86

Sign Extend Example



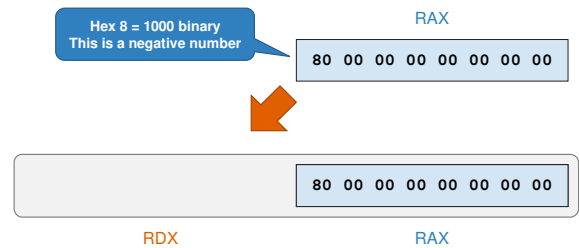
Fall 2024

Secureware State - Cook - CSU 35

87

87

Sign Extend Example



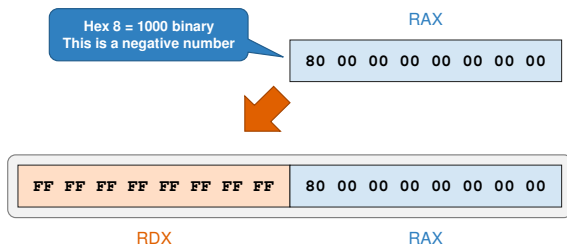
Fall 2024

Secureware State - Cook - CSU 35

88

88

Sign Extend Example



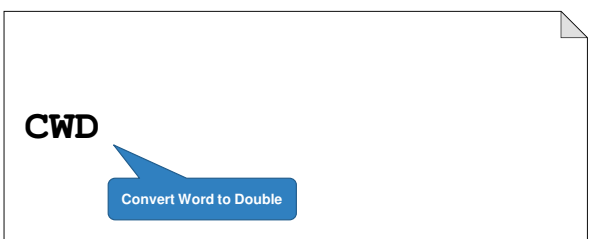
Fall 2024

Secureware State - Cook - CSU 35

89

89

CWD (16 bit): Extend AX → DX



Fall 2024

Secureware State - Cook - CSU 35

90

90

CDQ (32 bit): Extend EAX → EDX

CDQ

Convert Double to Quad

91

CQO (64 bit): Extend RAX → RDX

CQO

Use this one!

Convert Quad to Oct

92

Divide 64-bit: -1846 by 42

```
MOV rax, -1846    #RAX is the dividend
MOV rbx, 42       #Divisor
CQO              #Sign extend to RDX
IDIV rbx          #RAX gets quotient
                 #RDX gets remainder
```

93



How Compare Works

It's all math

94

Behind the scenes...



- The second argument is subtracted from the first
- The result of this computation is used to determine how the operands compare
- This subtraction result is discarded

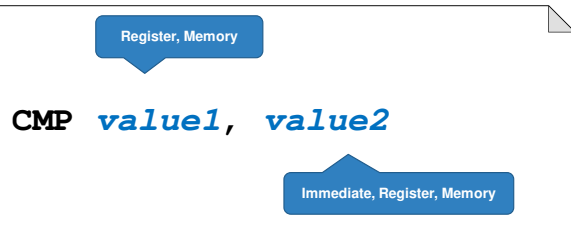
95

But... why subtract?

- Why subtract the operands?
- The result can tell you which is larger
- For example: A and B are both positive...
 - $A - B \rightarrow$ positive number \rightarrow A was larger
 - $A - B \rightarrow$ negative number \rightarrow B was larger
 - $A - B \rightarrow$ zero \rightarrow both numbers are equal

96

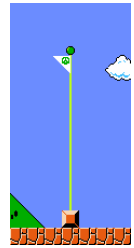
Instruction: Compare



97

Flags

- A *flag* is a Boolean value that indicates the result of an action
- These are set by various actions such as calculations, comparisons, etc...



98

Flags

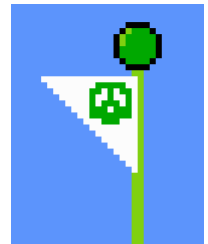
- Flags are typically stored as individual bits in the *Status Register*
- You can't change the register directly, but numerous instructions use it for control and logic



99

Zero Flag (ZF)

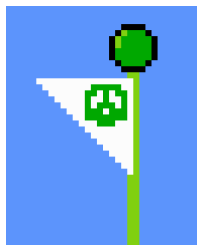
- True if the last computation resulted in zero (all bits are 0)
- For compare, the zero flag indicates the two operands are equal
- Used by quite a few conditional jump statements



100

Sign Flag (SF)

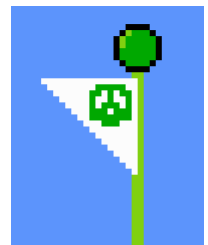
- True if the *most significant bit* of the result is 1
- This would indicate a *negative* 2's complement number
- Meaningless if the operands are interpreted as unsigned



101

Carry Flag (CF)

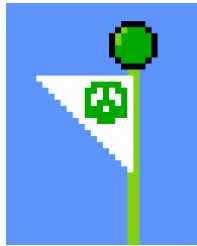
- True if a 1 is "borrowed" when subtraction is performed
- ...or a 1 is "carried" from addition
- For *unsigned* numbers, it indicates:
 - exceeded the size of the register on addition
 - or an underflow (too small value) on subtraction



102

Overflow Flag (OF)

- Also known as "signed carry flag"
- True if the sign bit changed *when it shouldn't have*
- For example:
 - (negative – positive) should be negative
 - a positive result will set the flag
- For signed numbers, it indicates:
 - exceeded the register size
 - i.e. the value was too big/small



Fall 2024

Secureworks, State - Cook - CSO 35

103

103

x86 Flags Used by Compare

Name	Description	When True
CF	Carry Flag	If a bit was "carried" or "borrowed" during math.
ZF	Zero Flag	All the bits in the result are zero.
SF	Sign Flag	If the most significant bit is 1.
OF	Overflow Flag	If the sign-bit changed when it shouldn't have.

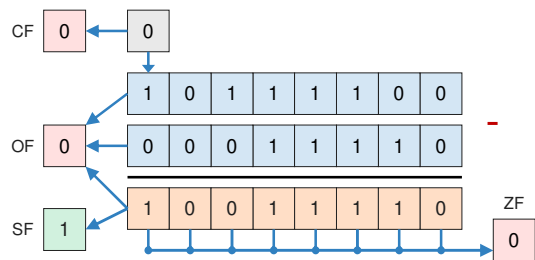
Fall 2024

Secureworks, State - Cook - CSO 35

104

104

-68 vs. 30 (if interpreted as signed)
188 vs. 30 (if interpreted as unsigned)



Fall 2024

Secureworks, State - Cook - CSO 35

105

105

Jump on Equality

Jump	Description	When True
JE	Equal	ZF = 1
JNE	Not equal	ZF = 0

Fall 2024

Secureworks, State - Cook - CSO 35

106

106

Signed Jump Instructions

Jump	Description	When True
JG	Jump Greater than	SF = OF, ZF = 0
JGE	Jump Greater than or Equal	SF = OF
JL	Jump Less than	SF ≠ OF, ZF = 0
JLE	Jump Less than or Equal	SF ≠ OF

Fall 2024

Secureworks, State - Cook - CSO 35

107

107

Unsigned Jumps

Jump	Description	When True
JA	Jump Above	CF = 0, ZF = 0
JAEC	Jump Above or Equal	CF = 0
JB	Jump Below	CF = 1, ZF = 0
JBEC	Jump Below or Equal	CF = 1

Fall 2024

Secureworks, State - Cook - CSO 35

108

108

Unsigned Conditional Jump Example

```
_start:  
    mov    rax, 42  
    cmp    rax, 13  
    jae    Bigger  
    ...  
Bigger:  
    add    rax, 5
```

rax >= 13?