



## Memory Indexing

Part 7

1



## Sizing Instructions

How many bytes are you using?

2

### Sizing Instructions

- The Intel can load/store 1-byte, 2-byte, 4-byte or 8-byte values
- Whenever a processor accesses memory, the instruction specifies how many bytes to access



Fall 2024

Section 7: Memory - CS: 10.10

3

3

### Sizing Instructions

- The assembler will automatically fill this information in for you.
- *How?* If a register is used, the assembly can assume it by looking at size of the register



Fall 2024

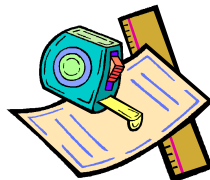
Section 7: Memory - CS: 10.10

4

4

### Sizing Instructions

- However, sometimes the number of bytes (1, 2, etc..) can't be determined
- In this case, the assembler will report an error
- ... since it doesn't know how to encode the instruction



Fall 2024

Section 7: Memory - CS: 10.10

5

5

### Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global Main
Main:
    mov total, 50
```

total is a target address.  
It doesn't have any  
implied size.

Fall 2024

Section 7: Memory - CS: 10.10

6

6

## Example: How Many Bytes?

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global Main
Main:
    mov total, 50
```

How many bytes is this?  
The value 50 can be  
stored in 1, 2, 4, or 8  
bytes.

Fall 2024

Secrets of the Shell - CS50.95

7

## How Many Bytes?

- If the assembler can't infer how many bytes to access, it'll report *"ambiguous operand size"*
- To address this issue...
  - GAS assembly allows you place a single character after the instruction's mnemonic
  - this suffix will tell the assembler how many bytes will be accessed during the operation

Fall 2024

Secrets of the Shell - CS50.95

8

## How Many Bytes

Suffix	Name	Size
<b>b</b>	byte	1 byte
<b>s</b>	short	2 bytes
<b>l</b>	long	4 bytes
<b>q</b>	quad	8 bytes

Fall 2024

Secrets of the Shell - CS50.95

9

## Example: Suffix Used

```
.intel_syntax noprefix
.data
total:
    .quad 0

.text
.global Main
Main:
    movq total, 50
```

Note the **q**.  
Now the assembler knows  
you mean "move quad".

Fall 2024

Secrets of the Shell - CS50.95

10

## Behind the Scenes of Arrays

All the mystery is revealed!

## Arrays

- Computers do not have an 'array' data type
- So, how do you have array variables?
- When you create an array...
  - you allocate a block of memory
  - each element is located sequentially in memory – one right after each other

Fall 2024

Secrets of the Shell - CS50.95

12

## Arrays

- Every byte in memory has an address
- This is just like an array
- To get an array element
  - we merely need to **compute** the address
  - we must also remember that some values take multiple bytes – **so there is math**

Fall 2024

Secrets of the Machine - CS50.10

13

13

## Array Math Example

- Let's again assume that our buffer starts at address **2000**
- The first array element is located at address 2000
- Arrays consists of bytes...
  - the second is **2001**
  - the third is **2002**
  - the fourth **2003**
  - etc...

2000	H
2001	e
2002	l
2003	l
2004	o

Fall 2024

Secrets of the Machine - CS50.10

14

14

## Array Math Example – 16 bit

- First element uses 2000... 2001
- Since each array element is 2 bytes...
  - second address is **2002**
  - third address is **2004**
  - fourth address is **2006**
  - etc...

2000	F0A3
2002	042B
2004	C1F1
2006	0D0B
2008	9C2A

Fall 2024

Secrets of the Machine - CS50.10

15

15

## Array Math Example – 64 bit

- First element uses 2000 to 2007
- Second address is **2008**
- Third address is **2016**
- Fourth address is **2024**
- etc...

2000	446576696E20436F
2008	6F6B000000000000
2016	53616372616D656E
2024	746F205374617465
2032	4353433335000000

Fall 2024

Secrets of the Machine - CS50.10

16

16

## Behind the Scenes...

- So, when an array element is read, internally, a mathematical equation is used
- It uses the start of the first element, the array index, and the size of each element

<code>start address + (index × size)</code>
---

Fall 2024

Secrets of the Machine - CS50.10

17

17

## Behind the Scenes...

- This is why the C Programming Languages uses zero as the first array element**
- If zero is used with this formula, it gets the start of the buffer

<code>start address + (index × size)</code>
---

Fall 2024

Secrets of the Machine - CS50.10

18

18

## Behind the Scenes...

- Java uses zero-indexing because C does
- ... and C does so it can create efficient assembly!

```
start address + (index × size)
```

19



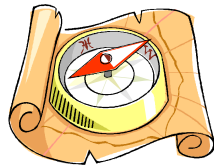
## Indexing on the x64

Grabbing any byte

20

## Indexing on the x64

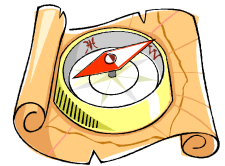
- The Intel x64 supports direct, indirect, indexing and scaling
- So, the Intel is very versatile in how it can access memory
- This is typical of CISC-ish architectures



21

## Effective Addresses

- Processors have the ability to create the *effective address* by combining data
- How it works:
  - starts with a base address
  - then adds a value (or values)
  - finally, uses this temporary value as the actual address



22

## Effective Addresses

- Using the addresses stored in memory, registers, etc... is useful in programs
- Often programs contain *groups* of data
  - fields in an abstract data type
  - elements in an array
  - entries in a large table etc...



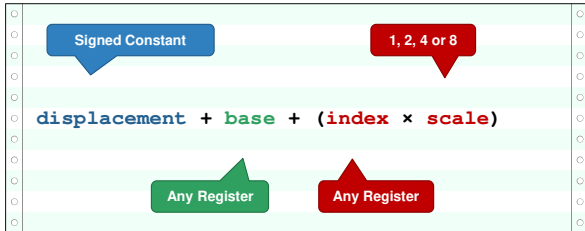
23

## Terminology

- *Base-address* is the initial address
- *Displacement (aka offset)* is a constant (immediate) that is added to the address
- *Index* is a *register* added to the address
- *Scale* used to multiply the index before adding it to the address

24

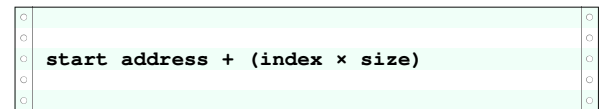
## x64 Effective Address Formula



25

## Behind the Scenes...

- But wait, doesn't that formula look familiar?
- The addressing term "scale" is basically equivalent to "size" in this example
- Addressing and arrays work together flawlessly



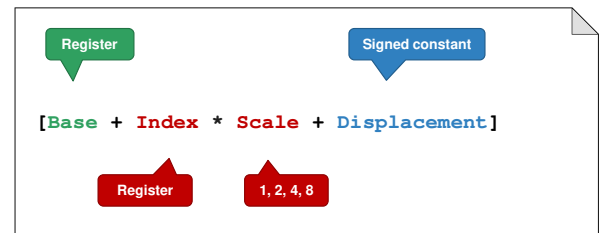
26

## Addressing Notation in Assembly

- Intel Notation (*Microsoft actually created it*) allows you to specify the full equation
- The notation is very straight forward and mimics the equation used to compute the effective address
- Parts of the equation can be omitted, and the assembler will understand

27

## Intel Notation



28

## Notation (reg = register)

Mode	Syntax	Java Equivalent
Immediate	value	value
Register	register	register
Direct	label	Memory[label]
Direct Indexed	[label + reg]	Memory[label + reg]
Indirect	[reg]	Memory[reg]
Indirect Indexed	[reg + reg]	Memory[reg + reg]
Indirect Indexed Scale	[reg + reg * scale]	Memory[reg + reg * scale]

29

## Addressing Notation in Assembly

- When you write an assembly instruction...
  - you specify all 4 four addressing features
  - however, notation fills in the "missing" items
- For example: for direct addressing...
  - Displacement → Address of the data
  - Base → Not used
  - Index → Not used
  - Scale → 1, irrelevant without an Index

30

## Indexing Examples

- The following examples use addressing modes to modify an ASCII buffer
- Let's **assume** that the start of the buffer **Talk** is **5000**

Talk = 5000

5000	48	H
5001	65	e
5002	6C	l
5003	6C	l
5004	6F	o

Fall 2024

Segment: Base - Cx86 - CS:32

31

31

## Example: Direct Index

```
# Talk = 5000
mov rdi, 1
movb [Talk + rdi], 33
```

Using the rsi register for indexing, but you can use any register

ASCII 33 → !

5000	48	H
5001	33	!
5002	6C	l
5003	6C	l
5004	6F	o

Fall 2024

Segment: Base - Cx86 - CS:32

32

32

## Example: Direct Index (Scale 2)

```
# Talk = 5000
mov rdi, 1
movb [Talk + rdi * 2], 33
```

5000	48	H
5001	65	e
5002	33	!
5003	6C	l
5004	6F	o

Fall 2024

Segment: Base - Cx86 - CS:32

33

33

## Example: Direct Index (Scale 4)

```
# Talk = 5000
mov rdi, 1
movb [Talk + rdi * 4], 33
```

5000	48	H
5001	65	e
5002	6C	l
5003	6C	l
5004	33	!

Fall 2024

Segment: Base - Cx86 - CS:32

34

34

## Example: Register Indirect

```
# Talk = 5000
lea rax, Talk
movb [rax], 33
```

Indirect. Base is rax

5000	33	!
5001	65	e
5002	6C	l
5003	6C	l
5004	6F	o

Fall 2024

Segment: Base - Cx86 - CS:32

35

35

## Example: Register Indirect Index

```
# Talk = 5000
lea rax, Talk
mov rdi, 1
movb [rax + rdi], 33
```

Base Index

5000	48	H
5001	33	!
5002	6C	l
5003	6C	l
5004	6F	o

Fall 2024

Segment: Base - Cx86 - CS:32

36

36

### Ex: Register Indirect Index (Scale 2)

```
# Talk = 5000

lea rax, Talk
mov rsi, 1
movb [rax + rsi * 2], 33
```

5000	48	H
5001	65	e
5002	33	!
5003	6C	l
5004	6F	o

Scale

37

### Ex: Register Indirect Index (Scale 4)

```
# Talk = 5000

lea rax, Talk
mov rsi, 1
movb [rax + rsi * 4], 33
```

5000	48	H
5001	65	e
5002	6C	l
5003	6C	l
5004	33	!

38



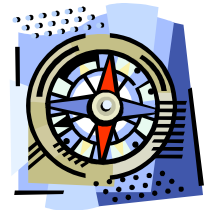
### Addressing & Loops

They were made for each other ... *literally*

39

### Addressing & Loops

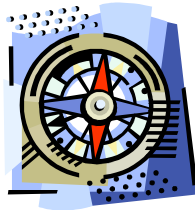
- When you use arrays in Java, often the index is a variable
- This allows you to use a For Loop to analyze very element in the array
- This is more common than you think in assembly



40

### Addressing & Loops

- So, processors allow a register to be used as an index
- This allows you to:
  - copy strings (copying arrays)
  - search through a list
  - and much more...



41

### For Loop: 0 to 4 - Before

```
.intel_syntax noprefix
.data
Greet:
    .ascii "HELLO"

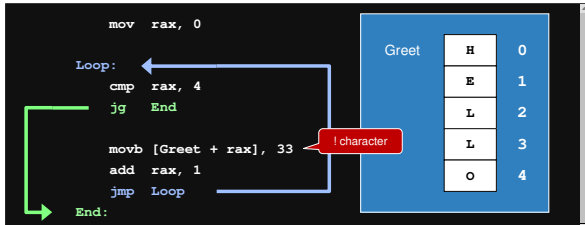
.text
.global Main

Main:
```

Greet	H	0
	E	1
	L	2
	L	3
	O	4

42

## For Loop: 0 to 4



43

## For Loop: 0 to 4 - After



44

## Tables

How to Organize Data

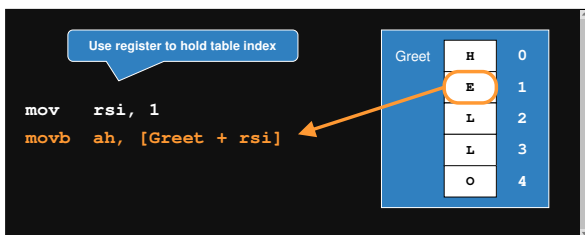
45

## Tables

- In assembly, you have full control of memory
- You can take advantage of these to create tables
- They can contain any data – from integers, to characters, to addresses

46

## Accessing Each element



47

## Tables of Integers

- Tables can contain *anything!*
- Often, they are used to store integers & addresses (8 bytes on a 64-bit system)
- Just make sure to use the scale feature!

48



## Table of Long Integers

Years:

```
.quad 1776
.quad 1783
.quad 1846
.quad 1850
.quad 1947
```

8 Bytes each

Fall 2024

SecureWeb: Stan - Cook - CSU 35

49

## Assuming Years is 6000

Years:

```
.quad 1776
.quad 1783
.quad 1846
.quad 1850
.quad 1947
```

6000	1776
6008	1783
6016	1846
6024	1850
6032	1947

Fall 2024

SecureWeb: Stan - Cook - CSU 35

50

## Assuming Years is 6000

Table index 1

```
mov rsi, 1
mov rcx, [Years + rsi * 8]
```

Note the scale!

6000	1776
6008	1783
6016	1846
6024	1850
6032	1947

Fall 2024

SecureWeb: Stan - Cook - CSU 35

51

## Table of Addresses. Assume Names is 3000

```
Sutter:
.ascii "John Sutter\0"

Marshal:
.ascii "James Marshal\0"

Names:
.quad Sutter
.quad Marshal
```

3000	Sutter (address)
3008	Marshal (address)

Fall 2024

SecureWeb: Stan - Cook - CSU 35

52

## Assuming Names is 3000

Note: mov is used. We want the data from the table (which is an address)

```
mov rsi, 1
mov rdx, [Names + rsi * 8]

call PrintString
```

3000	Sutter (address)
3008	Marshal (address)

Fall 2024

SecureWeb: Stan - Cook - CSU 35

53



## Buffer Overflow

With Great Power  
Comes Great Responsibility

54

## Buffer Overflow

- Operating systems protect programs from having their memory / code damaged by *other* programs
- However...operating systems don't protect programs from damaging *themselves*



Fall 2024

Secureworks State - Cisk - CSU 35

55

55

## Buffers & Programs

- In memory, a running program's data is often stored next to its instructions
- This means...
  - if the end of a buffer of exceeded, the program can be read/written
  - this is a common hacker technique to modify a program *while it is running!*

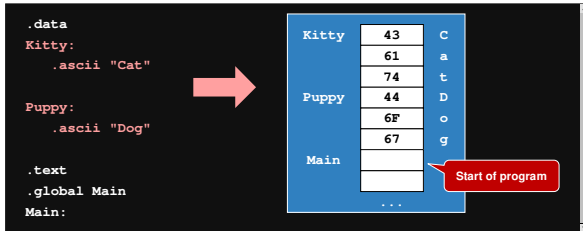
Fall 2024

Secureworks State - Cisk - CSU 35

56

56

## Example Program



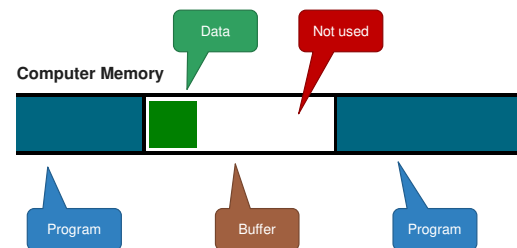
Fall 2024

Secureworks State - Cisk - CSU 35

57

57

## Buffer Overflow – How it Works



Fall 2024

Secureworks State - Cisk - CSU 35

58

58

## Buffer Overflow



- It is possible to store too much information – resulting in a *buffer overflow*
- The extra bytes will overwrite part of the running program – changing it!

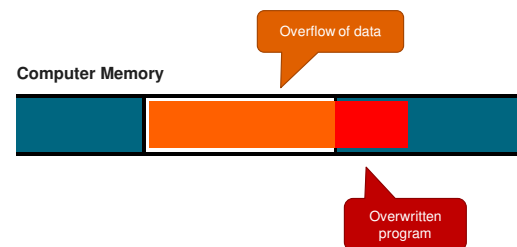
Fall 2024

Secureworks State - Cisk - CSU 35

59

59

## Buffer Overflow – How it Works



Fall 2024

Secureworks State - Cisk - CSU 35

60

60

## Bad Indexing

- It is possible to accidentally change data stored in the different buffers
- In assembly, you have full control over your allocated memory
- With great power comes great responsibility*



Fall 2024

SecureCode - CS61B

61

61

## Wrong Buffer Changed

```
.intel_syntax noprefix
.data
Kitty:
    .ascii "Cat\0"
Puppy:
    .ascii "Dog\0"

.text
.global Main
Main:
    mov rdi, 4
    movb [Kitty + rdi], 72
```

4 bytes. Character indexes from 0 to 3

72 is ASCII 'H'  
In hex it's 48

Fall 2024

SecureCode - CS61B

62

62

## Wrong Buffer Changed

```
.intel_syntax noprefix
.data
Kitty:
    .ascii "Cat\0"
Puppy:
    .ascii "Dog\0"

.text
.global Main
Main:
    mov rdi, 4
    movb [Kitty + rdi], 72
```



Kitty	43	C
	61	a
	74	t
	00	
Puppy	44	D
	6F	o
	67	g
	00	

Fall 2024

SecureCode - CS61B

63

63

## Wrong Buffer Changed

```
.intel_syntax noprefix
.data
Kitty:
    .ascii "Cat\0"
Puppy:
    .ascii "Dog\0"

.text
.global Main
Main:
    mov rdi, 4
    movb [Kitty + rdi], 72
```

Kitty	43	C
	61	a
	74	t
	00	
Puppy	48	H
	6F	o
	67	g
	00	

Fall 2024

SecureCode - CS61B

64

64

## Endianness



The "proper" order of things

## So Many Bytes...

- On a 64-bit system, each word consists of 8 bytes
- So, when any 64-bit value is stored in memory, each of those 8 bytes must be stored
- However, question remains: *What order do we store them?*



Fall 2024

SecureCode - CS61B

66

66

## Example Unsigned Integer (4 Byte)

1,188,852,977

46	DC	74	F1
----	----	----	----

Most significant Byte (MSB)

Least significant Byte (LSB)

Fall 2024

Secureworks, State - Cook - CSU 35

67

## So Many Bytes...

- Do we store the least-significant byte (LSB) first, or the most-significant (MSB)?
- As long as a system always follows the same format, then there are no problems
- ... but different system use different approaches

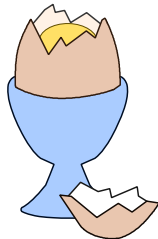
Fall 2024

Secureworks, State - Cook - CSU 35

68

## Big Endian vs. Little Endian

- Big-Endian approach
  - store the MSB first
  - used by Motorola & PowerPC
- Little-Endian approach
  - store the LSB first
  - used by Intel



Fall 2024

Secureworks, State - Cook - CSU 35

69

## Big Endian vs. Little Endian

46	DC	74	F1
----	----	----	----

Big Endian	
0	46
1	DC
2	74
3	F1

Little Endian	
0	F1
1	74
2	DC
3	46

Fall 2024

Secureworks, State - Cook - CSU 35

70

## Assuming Value is located at 2000

Value:  
.quad 74

2000	4A
2001	00
2002	00
2003	00
2004	00
2005	00
2006	00
2007	00

Least Significant Byte (LSB)

Little Endian

Fall 2024

Secureworks, State - Cook - CSU 35

71

## No "End" to Problems

- There is a problem...*  
if two systems use different formats, data will be interpreted incorrectly!
- If how the read differs from how it is stored, the data will be mangled



Fall 2024

Secureworks, State - Cook - CSU 35

72

## No "End" to Problems

- For example:
  - a **little**-endian system reads a value stored in **big**-endian
  - a **big**-endian system reads a value stored in **little**-endian
- Programmers must be conscience of this whenever binary data is accessed



Fall 2024

Secureworks, State - Cook - CSU 35

73

73

## No "End" to Problems

- So, whenever data is read from secondary storage, you **cannot** assume it will be in your processor's format
- This is compounded by file formats (gif, jpeg, mp3, etc...) which are also inconsistent



Fall 2024

Secureworks, State - Cook - CSU 35

74

74

## Example File Format Endianness

File Format	Endianness
Adobe Photoshop	Big Endian
Windows Bitmap (.bmp)	Little Endian
GIF	Little Endian
JPEG	Big Endian
MP4	Big Endian
ZIP file	Little Endian

Fall 2024

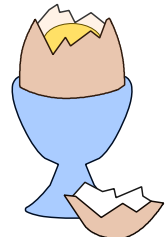
Secureworks, State - Cook - CSU 35

75

75

## So... who is correct?

- So, what is the correct and superior format?
- Is it Intel (little endian)?
- ...or the PowerPC (big endian) correct?



Fall 2024

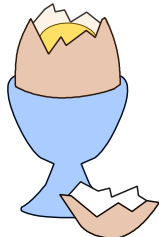
Secureworks, State - Cook - CSU 35

76

76

## So... who is correct?

- In reality neither side is superior
- Both formats are equally correct
- Both have minor advantages in assembly... but nothing huge



Fall 2024

Secureworks, State - Cook - CSU 35

77

77

## Gulliver's Travels



Fall 2024

Secureworks, State - Cook - CSU 35

78

78