




Subroutines & Operating Systems

Part 8

1




The System Stack

Pile of... Data

2

The System Stack

- The processor maintains a stack in memory
- It allows *subroutines*
 - analogous to the "functions" you use in Java and other third-generation languages
 - but, much more simple



Fall 2024

Subroutines: Stack - CS61B

3

3

Examples of Stacks

- Page-visited "back button" history in a web browser
- Undo sequence in a text editor
- Deck of cards in Windows Solitaire



Fall 2024

Subroutines: Stack - CS61B

4

4

Implementing in Memory

- On a processor, the stack stores integers
 - size of the integer the bit-size of the system
 - 64-bit system → 64-bit integer
- Stacks is stored in memory
 - A fixed location pointer (S0) defines the bottom of the stack
 - A *stack pointer* (SP) gives the location of the top of the stack

Fall 2024

Subroutines: Stack - CS61B

5

5

Approaches

- Growing upwards
 - Bottom Pointer (S0) is the *lowest* address in the stack buffer
 - stack grows towards *higher* addresses
- Grow downwards
 - Bottom Pointer (S0) is the *highest* address in the stack buffer
 - stack grows towards *lower* addresses

Fall 2024

Subroutines: Stack - CS61B

6

6

Size of the Stack

- As an abstract data structure...
 - stacks are assumed to be **infinitely** deep
 - so, an arbitrary amount of data can be stored
- However...
 - stacks are implemented using memory buffers
 - which are **finite** in size
- If the data **exceeds** the allocated space, a **stack overflow** error occurs

Fall 2024

Secureware: Stack - Crash - CS50.10

7

7

Subroutine Call Basics



Organizing Your Program

8

Subroutine Call

- The stack is essential for subroutines to work
- How?
 - used to save the return addresses for call instructions
 - backup and restore registers
 - pass data between subroutines



Fall 2024

Secureware: Stack - Crash - CS50.10

9

9

When you call a subroutine...

1. Processor pushes the instruction pointer (IP) – an address – on the stack
2. IP is set to the address of the subroutine
3. Subroutine executes and ends with a "return" instruction
4. Processor pops & restores the original IP
5. Execution continues after the initial call

Fall 2024


Secureware: Stack - Crash - CS50.10

10

10

Nesting is Possible

- Subroutines can call other subroutines
- f()** calls **g()** which then calls **h()**, etc...
- The stack stores the return addresses of the callers
- Just like the "history button" in your web browser, you can store many return addresses



Fall 2024


Secureware: Stack - Crash - CS50.10

11

11

Nesting is Possible

- Each time a subroutine completes, the processor pops the top of the stack
- ...then returns to the **caller**
- This allows normal function calls and recursion (a powerful tool)




Fall 2024

Secureware: Stack - Crash - CS50.10

12

12




x64 Subroutines

Organizing Your Programs ... with Intel

13

Instruction: Call

- The *Call Instruction* transfers control to a subroutine
- Other processors call it different names such as JSR (Jump Subroutine)
- The stack is used to save the current IP



Fall 2024 Sacramento State - CSIS - CSIS 35 14

14

Instruction: Call

CALL *address*

Usually, a label
(which is an address)

Fall 2024 Sacramento State - CSIS - CSIS 35 15

15

Instruction: Return

- The Return Instruction is used mark the end of subroutine
- When the instruction is executed...
 - the old instruction pointer is read from the system stack
 - the current instruction pointer is updated – restoring execution after the initial call

Fall 2024 Sacramento State - CSIS - CSIS 35 16

16

Instruction: Return

- Do not forget this!
- If you do...
 - execution will simply continue, in memory, until a return instruction is encountered
 - often is can run past the end of your program
 - ...and run data!

Fall 2024 Sacramento State - CSIS - CSIS 35 17

17

Instruction: Return

RET

No arguments!

Fall 2024 Sacramento State - CSIS - CSIS 35 18

18

Subroutine Example

```

Main:
    mov rcx, 4
    mov rbx, 12
    call AddIt
    add rbx, 1
    ...
AddIt:
    add rbx, rcx
    ret
    
```

19



Saving Registers & Lost Data

Making subroutines clean

20

Saving Registers & Lost Data

- Each subroutine will use the registers as it needs
- So, when a subroutine is called, *it may modify the caller's registers*



21

Subroutine Example

```

StingersUp:
    .ascii "Stingers Up!\n\0"
    ...
Main:
    mov rdx, 1947
    call Hornet
    call PrintInteger
    
```

22

```

Main:
    mov rdx, 1947
    call Hornet
    call PrintInteger
    ...
Hornet:
    lea rdx, StingersUp
    call PrintString
    ret
    
```

23

Subroutine Wrong Output

Stingers Up!
4202698

24

Saving Registers & Lost Data



- Some processors have few registers – this is very likely
- This can lead to hard-to-fix bugs if caution is not used – *e.g. subroutine changes the callers loop counter*

Fall 2024


Securely Store - Cook - CSU 35

25

25

Use the Stack!

- Subroutine saves registers it will change
- How? Registers are pushed onto the state at the beginning of the subroutine
- Before it returns, it pops (and restores) the old values



Fall 2024

Securely Store - Cook - CSU 35

26

26

Saving Registers... How nice! :-)

DoSomething:

```
push rax
push rbx
push rcx
```

Backup registers

```
...
```

Your code

```
pop rcx
pop rbx
pop rax
ret
```

Restore them.
Note the reverse order

Fall 2024

Securely Store - Cook - CSU 35

27

27

Fixed Subroutine

Hornet:

```
push rdx
```

Push the value on the stack
stack.push(rdx)

```
lea rdx, StingersUp
call PrintString
```

```
pop rdx
ret
```

Restore the value
rdx = stack.pop()

Fall 2024

Securely Store - Cook - CSU 35

28

28

Subroutine Correct Output

```
Stingers Up!
1947
```

The subroutine
worked perfectly

Fall 2024

Securely Store - Cook - CSU 35

29

29




Stack Frames

The Jenga of data!

30

Stack Frames

- *Stack Frames* is an compiler approach where subroutine arguments are passed using the system stack
- The stack is also used to store local variables



Fall 2024


Subroutines: Stack - C&C++ CS 31

31

31

Why is this needed?

1. Need to support **any** number of parameters – *even if it exceeds the available number of registers*
2. Need support local variables



Fall 2024

Subroutines: Stack - C&C++ CS 31

32

32

Stack Frame Contents

- Contains all the information needed by subroutine
- Includes:
 - calling program's return address
 - input parameters to the subroutine
 - the subroutine's local variables
 - space to backup the caller's register file

Fall 2024


Subroutines: Stack - C&C++ CS 31

33

33

Nesting is Possible

- Stack is LIFO (last in first out), so subroutines can call subroutines
- This approach allows recursion and all the features found in high-level programming languages



Fall 2024


Subroutines: Stack - C&C++ CS 31

34

34

Nesting is Possible

- For example, subroutine **f()** calls **g()** which then calls **h()**
- Since the stack is used, the stack frames follow the LIFO behavior



Fall 2024

Subroutines: Stack - C&C++ CS 31

35

35

How it Starts Up

- Caller
 - pushes the subroutine's arguments onto the stack
 - caller calls the subroutine
- Subroutine then...
 - uses the stack to backup registers
 - and *"carve"* out local variables

Fall 2024

Subroutines: Stack - C&C++ CS 31

36

36

How it Finishes

- Subroutine...
 - restores the original register values
 - removes the local variables from the stack
 - calls the processor "return" instruction
- Caller, then...
 - removes its arguments from the stack
 - handles the result – which can be passed either in a register or on the stack


Fall 2024

Segment 8: Stack - CSU 35

37

37

Stack Frame Size Varies



- The number of input arguments and local variables varies from subroutine to subroutine
- The arrangement of data within the stack frame also varies from compiler to compiler
- Stack frames is a *concept* – and it is used with various differences

Fall 2024

Segment 8: Stack - CSU 35

38

38

What About Different Object Files?

- Programs are often created from multiple object files
- These can be created by different compilers and linked *separately* –
- So, how do we make sure that these are all compatible?



Fall 2024

Segment 8: Stack - CSU 35

39

39

Calling Convention

- A *calling convention* is defined by a programming system (e.g. a language) to define *how* data will be passed
- In particular, it defines the structure of the stack frame and how data is returned



Fall 2024

Segment 8: Stack - CSU 35

40

40

Calling convention

- For example:
 - Is the first argument pushed first? Or last?
 - Is the result in a register? Or the stack?
- If all subroutines follow the same format
 - caller can use the same format for each
 - subroutines can also be created separately and linked together

Fall 2024

Segment 8: Stack - CSU 35

41

41

Compatibility

- If two different compilers use the same calling convention, the resulting object files will be compatible
- This means, large applications can be created in different programming languages



Fall 2024

Segment 8: Stack - CSU 35

42

42




Stack Frames on the x64

Pretty much how its done on all processors

43

Stack Frames on the x64


- Stack frames on the x64 are accomplished pretty much the same way as other processors
- How it is done in real life is not simple – and is one of the hardest concepts to understand



44

Stack Frames on the x64

- On the x64, we will use the Base Pointer (RBP) to access elements in the stack frame
- This is a pointer register
- We will use it as an "anchor" in our stack frame



45

Stack Frames on the x64

- As we build the stack frame, we will set RBP to fixed address in the stack frame
- Our parameters and local variables will be accessed by looking at memory *relative* to the RBP
- So, we will look *x* many bytes above and below the "anchor"

46

Stack on the x64


- The stack base on the x64 is stored in high memory and grows downwards towards 0
- So, as the size of the stack increases, the stack pointer (RSP) will *decrease* in value



47

Stack on the x64

- On a 64-bit system, it will decrease by increments of 8 bytes
- So, each of our values (local variables and parameters) will be offsets of 8



48

Stack Frame Start Steps

1. Caller pushes arguments
2. Call the subroutine
3. Backup (push) the Base Pointer
4. Set the Base Pointer
5. "Carve" Local Variables
6. Backup (push) local variables



Fall 2024


Secureware State - CSIS - CSIS 35

49

49

Structure of a Stack Frame

- In this example, the stack is growing downward
- The stack pointer (RSP) is always on the bottom of the stack frame



Fall 2024


Secureware State - CSIS - CSIS 35

50

50

1. Push Parameters

- Caller starts by pushing each of the parameters onto the stack
- Order parameters are pushed is defined in the *calling convention*



Fall 2024


Secureware State - CSIS - CSIS 35

51

51

2. Call the subroutine

- The caller then uses the *Call Instruction* to pass control to the subroutine
- The processor pushes the IP (instruction pointer) on the stack
- Subroutine now runs



Fall 2024


Secureware State - CSIS - CSIS 35

52

52

3. Backup the Old Base Pointer

- We need to set the Base Pointer (RBP)
- So, the old version needs to be saved (so it can be restored)
- Old RBP is pushed



Fall 2024


Secureware State - CSIS - CSIS 35

53

53

4. Set the Base Pointer

- Then, it sets the Base Pointer (RBP) to the current stack pointer (RSP) address
- RBP is an "anchor" that we will use



Fall 2024


Secureware State - CSIS - CSIS 35

54

54

5. "Carve" Local Variables

- Subroutine now creates local variables on the stack
- Their initial values can be simply pushed



Fall 2024


Secureware State - Cook - CSU 35

55

55

6. Backup Registers

- Finally, the subroutine saves all the registers (that will change) on the stack
- It will restore them at the end



Fall 2024


Secureware State - Cook - CSU 35

56

56

The Completed Stack Frame

- The Stack Frame contains all the information a subroutine needs
- We just need to be able to access the data programmatically



Fall 2024


Secureware State - Cook - CSU 35

57

57

Parameters & Local Variables

- Now, RBP is set to an address between the parameters and the local variables
- We can use *offsets* from RBP to access each



Fall 2024

Secureware State - Cook - CSU 35

58

58

Compiler Friendly... Not Human



- The offset values have to count the *bytes* from the RBP register
- These are not easy to read since both parameters and local variables are just offsets

Fall 2024


Secureware State - Cook - CSU 35

59

59

Size of Stack Values – 64 bit

- x64 stack grows downward, so...
- Local variables will have a *negative* offset
- Parameters will have a *positive* offset



Fall 2024


Secureware State - Cook - CSU 35

60

60

Size of Stack Values – 32 bit


- On a 32-bit system, each word is 4 bytes
- So, each value on the stack is 4 bytes
- Offsets increase and decrease by 4



61

Size of Stack Values – 64 bit

- On a 64-bit system, each word is 8 bytes
- So, each value on the stack is 8 bytes
- Offsets increase and decrease by 8



62

Caller Example

```
push 42
push rax
call Subroutine
```

1. Push parameters (16 bytes added)

2. Call subroutine

63

Subroutine: Setup Example

```
push rbp
mov rbp, rsp
push 1
push 2
push rax
push rbx
```

3. Backup RBP

3. Set new RBP

4. Local variables

5. Backup registers

64

Subroutine Data Example

```
mov rax, [rbp + 16]
add rax, [rbp + 24]
mov [rbp - 8], rax
```

Offset from pointer

Parameter 1

Parameter 2

Local variable 1

65

Stack Frame Ending Steps

- Restore registers (pop)
- Restore the stack pointer (set it to the base pointer)
- Restore the old base pointer
- Return and delete parameters



66

Stack Frame Return

- The Return can also be used to clean up the caller's stack items
- You can specify the number of bytes to pop (and discard) after the return
- Alternatively, the caller can clean up the stack

Fall 2024

Secureware State - Cook - CSU 35

67

67

Stack Frame Return

RET *byteCount*

Bytes to discard
after return

Fall 2024

Secureware State - Cook - CSU 35

68

68

Subroutine: Ending Example

```
pop    rbx
pop    rax

mov    rsp, rbp
pop    rbp
ret    16
```

1. Restore registers (pop in reverse order)

2. Set RSP to RBP
Effectively deletes all local variables

3. Restore RBP

4 Return after the stack is restored

Fall 2024

Secureware State - Cook - CSU 35

69

69

MySub :

```
push    rbp
mov     rbp, rsp
push    1
push    2
push    rax
push    rbx
```

Setup base pointer

Local variables (with initial values)

backup registers

...

```
pop     rbx
pop     rax
mov     rsp, rbp
pop     rbp
ret     16
```

Restore registers

Restore base pointer

Fall 2024

Secureware State - Cook - CSU 35

70

70



Operating Systems

The master software

What is an operating system?

- The operating system is simply a series of programs
- These programs, however, run with special privileges which are needed by the OS
- Processors support two modes for executing programs



Fall 2024


Secureware State - Cook - CSU 35

72

72

Execution Modes

- *Privileged (supervisor) mode*
 - can run special instructions
 - can talk to all the hardware
 - etc...
- *User mode*
 - can only execute certain instructions
 - can't talk to all the hardware



Fall 2024


Secureworks State - Cook - CSIS 35

73

73

Vector Tables

- Programs (and hardware) often need to talk to the operating system
- Examples:
 - software needs talk to the OS
 - USB port notifies the OS that a device was plugged in



Fall 2024


Secureworks State - Cook - CSIS 35

74

74

Vector Tables

- But how does this happen?
- The processor can be *interrupted* – *alerted* – that something must be handled
- It then runs a special program that handles the event



Fall 2024

Secureworks State - Cook - CSIS 35

75

75

Vector Table

- During an interrupt, the device sends the processor an *interrupt number*
- The processor looks up the number in the *vector table*
- Table contains the address of *Interrupt Service Routine (ISR)* to execute




Fall 2024

Secureworks State - Cook - CSIS 35

76

76

How It Works



Fall 2024

Secureworks State - Cook - CSIS 35

77

77

What Happens...

1. Device interrupts the processor
2. Current program state is saved
 - register file
 - instruction pointer
3. Processor executes ISR using the Vector Table address
4. Current program state is restored



Fall 2024

Secureworks State - Cook - CSIS 35

78

78

The Kernal

- All these Interrupt Service Routines belong to the *kernal* – the core of the operating system
- Vast majority of the operating system is hidden from the end user



Fall 2024

Secureworks, Inc. - Conf - CSO 20

79

79

Interact with Applications



How do WE talk to the OS

80

Interact with Applications

- Software also needs to talk to the operating system
- For example:
 - draw a button
 - print a document
 - close this program
 - etc...



Fall 2024

Secureworks, Inc. - Conf - CSO 20

81

81

Interact with Applications

- Software can interrupt itself with a specific number
- This interrupt is *designated specifically for software*
- The operating system then handles the software's request



Fall 2024

Secureworks, Inc. - Conf - CSO 20

82

82

Application Program Interface

- Programs "talk" to the OS using *Application Program Interface (API)*
- Application → Operating System → IO
- Benefits:
 - makes applications faster and smaller
 - also makes the system more secure since apps do not directly talk to IO

Fall 2024

Secureworks, Inc. - Conf - CSO 20

83

83

Instruction: syscall (64-bit)

SYSCALL

Calls interrupt number reserved for programs needing attention

Fall 2024

Secureworks, Inc. - Conf - CSO 20


84

84

Subroutine vs. Interrupt

Subroutine	Interrupt
Executes code	Executes code
Returns when complete	Returns when complete
Called by the application	Executed by the processor
Part of the application	Handles events for the OS

85




Linux System Calls

How software and hardware "talk"

86

Interrupts on the Linux


- Linux, like other operating systems communicate with applications using *interrupts*
- Applications do not know where (in memory) to contact the kernel – so they ask the processor to do it



87


How It Works

1. Fill the registers
2. Interrupt using *syscall* (or INT 0x80 if on 32-bit)
3. Any results will be stored in the registers



88


How to Call Linux – 64 bit



- The **rax** register must contain the *system call number*
- This number indicates what you asking the OS to do
- There are only **329** total calls in the entire 64-bit UNIX operating system!

89

How to Call Linux – 64 bit




- Different registers are used to hold data
- The order is also quite odd:
rdi, rsi, rdx, r10, r8

90

Kernels are Simple!

- Linux only has **1** write and **1** read system call
- The location, number of bytes, and device only change
"write x many bytes from address y to device z"
- So, writing to the screen, a file, a port, etc...use the same call!



Fall 2024

Systemic State - Cook - CSU 35

91

91

Some Linux 64 Calls

System Call	rax	rdi	rsi	rdx
read	0	file descriptor	address	max bytes
write	1	file descriptor	address	count
open	2	address	flags	mode
close	3	file descriptor		
get pid	39			
exit	60	error code		

Fall 2024

Systemic State - Cook - CSU 35

92

92

Linux 64: Sys Write

```
mov rax, 1
mov rdi, 1
lea rsi, address
mov rdx, length
syscall
```

Linux command for WRITE

1 = Screen

Call Linux

Fall 2024

Systemic State - Cook - CSU 35

93

93

Linux 64: Sys Read

```
mov rax, 0
mov rdi, 0
lea rsi, address
mov rdx, maxBytes
syscall
```

Linux command for READ

0 = Keyboard

Maximum number of bytes to read

Call Linux

Fall 2024

Systemic State - Cook - CSU 35

94

94

Write Example

```
SacState:
.ascii "Stinger's up!\n"  #\n counts as 1 character
...
mov rax, 1      #1 = write
mov rdi, 1      #1 = screen
lea rsi, SacState
mov rdx, 14     #14 bytes
syscall
```

Fall 2024

Systemic State - Cook - CSU 35

95

95